

TraitHorizon: Scalable Exploration of Large Image-Feature Paired Datasets

Jackson Jacobs^{1*}, Fan Fan^{1*}, Laura Barisoni^{2,3}, and Andrew Janowczyk^{1,4,5}

¹ Department of Biomedical Engineering, Emory University and Georgia Institute of Technology, Atlanta, Georgia, USA ² Department of Pathology, Division of AI and Computational Pathology, Duke University, Durham, North Carolina, USA ³ Department of Medicine, Division of Nephrology, Duke University, Durham, North Carolina, USA ⁴ Division of Precision Oncology, Department of Oncology, University Hospital of Geneva, Geneva, Switzerland ⁵ Division of Clinical Pathology, Department of Diagnostics, University Hospital of Geneva, Geneva, Switzerland ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a

Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

1. Summary

Collections of Objects of Interest (OOIs)—such as cells, tubules, or tissue patches—paired with high-dimensional, computationally derived feature vectors (e.g., morphology, texture, or deep-learned embeddings) are increasingly being generated in image-based biomedical research. These OOI-feature datasets are routinely explored to detect outliers for quality control, uncover patterns, and generate biological insights. However, as these datasets grow in size and complexity, researchers face several bottlenecks: repeated manual creation of static visualizations for each subpopulation, difficulties in efficiently plotting large numbers of high-dimensional feature vectors, and limited tooling for tracing individual feature values back to their source objects. To address these challenges, we developed TraitHorizon, a browser-based visualization platform for interactive exploration of large OOI-feature pair datasets. TraitHorizon integrates three synchronized visualization components: (a) a parallel coordinates plot for feature-vector-level exploration, (b) dynamic violin plots for per-feature distribution analysis, and (c) a tabular data grid linking each feature vector to its corresponding object. Its interactive interface enables real-time, scalable visualization of datasets containing hundreds of thousands of OOI-feature pairs. We demonstrate TraitHorizon's utility through a case study involving quality control of a digital pathology dataset comprising 260,201 segmented tubules in kidney biopsies, each characterized by 99 features, illustrating how the platform facilitates rapid, interpretable, and reproducible data interrogation.

TraitHorizon is publicly available for use and modification at <https://traithorizon.com/>.

2. Statement of Need

Research in image-based fields (e.g., computational pathology) often aims to discover trends in Objects of Interest (OOIs), associated with diagnosis, prognosis, and therapy response. Objects can consist of individual items (e.g., cells), higher-order structures (e.g., tubules or infiltration patterns), or even small image patches. From these OOIs, sets of numerical features termed feature vectors are often extracted to encode their attributes. Features can range from simple (e.g., area, stain intensity, texture) to more complex, hand-crafted (e.g., aspect ratio of peritubular capillaries (Yijiang Chen et al., 2025)), or even deep learning-derived descriptors (Ambekar et al., 2025; Y. Chen et al., 2023; Yijiang Chen et al., 2025; Echle et al., 2021; Fan Fan et al., 2025). These features can then be examined to assess their associations with

object phenotypes—such as biological signals (e.g., disease state) or image quality factors (e.g., blurriness)—or how they positively/negatively correlate with one another.

Such analysis often involve generating static, non-interactive, visualizations, such as a parallel coordinates plots (Heinrich & Weiskopf, 2013) or scatter plots with UMAP (McInnes et al., 2020) or t-SNE (Maaten & Hinton, 2008). Aided by first-order summary statistics (e.g., mean, standard deviation), these visualizations are inspected to uncover structure and outliers within the feature distribution. Multiple data-driven subpopulations often emerge from this initial exploration, prompting repeated generation of visualizations and summary metrics for each subpopulation. Quantitative and qualitative patterns uncovered within subpopulations can, in turn, guide comparative analyses—for example, assessing how a given feature or pattern differs between “diseased” and “normal” groups.

However, to date, this investigatory workflow remains time-intensive due to three main limiting factors:

1. **Lack of connection between an object’s image, feature values, and cohort-level visualization.** Analysts often discover outliers or clusters within cohort-level visualizations, prompting the need to trace them back to the associated images for visual inspection. Static plots can at best display only a modest number of images, forcing analysts to manually map plotted points to objects, incurring significant time costs. The ability to instantly view the object associated with a plotted feature vector is essential for rapid validation, interpretation, and insight generation.
2. **Exploring each subpopulation incurs high time cost.** Biomedical imaging datasets often include heterogeneous subpopulations—such as tissue types, disease subtypes, or experimental conditions—that warrant the generation of separate visualizations for detailed inspection. Manually plotting each subpopulation becomes inefficient in repetitive workflows, while programmatic (e.g., loop-based) plot generation requires subpopulations of interest to be identified and well-defined *a priori*. Dynamic filtering and responsive plots are necessary for users to iteratively explore subpopulations without incurring substantial time costs from repeated plot generation and review. Furthermore, the conditions/filters used to identify them should be easy to document, save, and reapply.
3. **Latency while rendering large datasets disrupts analysis.** OOI-feature datasets can easily reach hundreds of thousands of data points. When subpopulations are not yet clearly defined, a “global” view of the entire dataset is required before smaller subpopulations can be explored. Browser-based tools not designed for plotting at this scale can suffer from heavy memory utilization, latency, and even unresponsiveness/browser crashes while rendering such data visualizations. A capable visualization tool must leverage rendering techniques that remain efficient and within browser limitations as the number of objects scales.

These limiting factors motivate the development of interactive data visualization tools that can dynamically generate visualizations as the user explores their data. However, existing tools do not fully address all three limiting factors simultaneously. TensorBoard Projector (TensorFlow Team, 2025) and HoloViews (Installation — HoloViews V1.21.0, n.d.) support interactive plotting of large multivariate datasets but do not provide for image viewing. Conversely, DendroMap (Bertucci et al., 2022) facilitates qualitative exploration of large sets of images but does not link back to their feature vectors. These tools also do not natively support dynamic filtering by feature value. HistoQC (Janowczyk et al., 2019) provides image-linked plots and metrics, but it is tailored specifically to quality control use cases in digital pathology and thus does not support arbitrary user-defined features vectors or images.

To address these gaps, we introduce TraitHorizon, a browser-based application for interactive exploration of images alongside their feature vectors. TraitHorizon (a) directly links plotted data points to their source object images, enabling rapid validation and interpretation, (b) supports dynamic, multi-dimensional filtering with support for saving and loading filter configurations, and (c) efficiently renders hundreds of thousands of data points in-browser. By doing so,

94 TraitHorizon facilitates efficient exploratory analysis in object-based research workflows.

95 3. Implementation

96 TraitHorizon is a Flask-based application ([Pallets/Flask, n.d.](#)), making it ideal for collaborative
 97 environments when hosted over a network. The front end is built with JavaScript, HTML5,
 98 and CSS, and leverages visualization libraries such as SlickGrid ([SlickGrid Home, n.d.](#)), Parallel
 99 Coordinates (parcoords ([Yun, 2025](#))), and D3.js ([D3 by Observable | The JavaScript Library for Bespoke Data Visualization, n.d.](#)) to power its interactive dashboard. Its command-line
 100 interface ingests a single directory of object image files (e.g., .png, .jpg, .svg, or any other
 101 browser-supported format (“[Image File Type and Format Guide - Media | MDN,](#)” 2025)) and
 102 a tab-separated values (TSV) file, with each row containing an image file path followed by
 103 a user-determined number of tab-separated feature columns. As such, any tool outputting
 104 tabular data is compatible with TraitHorizon’s simple input format (e.g., HistoQC ([Janowczyk
 105 et al., 2019](#)) and CohortFinder ([Fan Fan et al., 2024](#))). TraitHorizon supports integer, float,
 106 scientific notation, and categorical features. Clickable URLs are also supported by an optional
 107 column titled “url”.

109 TraitHorizon is designed to meet the demands of datasets containing hundreds of thousands of
 110 objects. To minimize browser lag during plotting, five application features were implemented:

- 111 1. Non-blocking progressive rendering occurs in the parallel coordinates plot, completing in
 112 seconds to minutes depending on dataset size. Importantly, the user interface remains
 113 responsive during this rendering period, allowing users to continue interacting with the
 114 parallel coordinates plot and other functional components without straining browser
 115 responsiveness.
- 116 2. Parallel coordinate filter configurations and associated filtered object IDs can be exported
 117 and imported in a lean JSON format, without requiring a re-render of the parallel
 118 coordinates plot.
- 119 3. The data grid supports both pagination and alternatively “infinite scrolling”. This web
 120 design technique enables new content to be dynamically loaded, replacing old content
 121 as the user scrolls down and creating the illusion of a never-ending page within a fixed
 122 memory footprint.
- 123 4. Object images are loaded on-demand to limit browser memory, network utilization, and
 124 CPU usage.
- 125 5. Violin plots are also computed on-demand when the user hovers over a parallel axis,
 126 avoiding CPU usage associated with refreshing violin plots each time the parallel
 127 coordinates plot is updated.

128 These design choices allow TraitHorizon to operate smoothly within the constraints of modern
 129 web browsers and modest consumer-grade hardware, regardless of dataset size.

130 4. Use Case – Interactive Exploration of Tubular Pathomic 131 Features via TraitHorizon

132 To illustrate TraitHorizon’s functionality, we provide the following detailed use case from
 133 a recent digital pathology image-based biomarker study ([F. Fan et al., 2024](#); [Fan Fan et
 134 al., 2025](#)). In our use case, we focused on data exploration of 260,201 segmented tubules
 135 ([Fogo et al., 2016](#)) and their associated features (for a total of 99 features per tubule) as
 136 computed from 254 Periodic Acid-Schiff (PAS) stained kidney biopsies from the Nephrotic
 137 Syndrome Study Network (NEPTUNE) ([Barisoni et al., 2013](#); [Gadegbeku et al., 2013](#)) and
 138 Cure Glomerulonephropathy (CureGN) ([CureGN Study Rationale, Design, and Methods, n.d.](#))
 139 consortia. The goal was to support quality control and uncover tubular morphological patterns
 140 along a pathway from normal to tubular atrophy ([Automated Computational Detection of](#)

141 [Interstitial Fibrosis, Tubular Atrophy, and Glomerulosclerosis | American Society of Nephrology,](#)
142 [n.d.; Fogo et al., 2016](#)).

143 4.1. Overview and Dataset Preparation

144 Before running TraitHorizon, users must prepare a TSV file containing the following columns:

- 145 1. **“filename” (required)** – Specifies the base filename for each instance image, without
146 the folder path (e.g., image1.png).
- 147 2. **Features (required)** – A tab separated row containing all features associated with each
148 instance. Our dataset included 99 tubular pathomic features capturing morphological and
149 topological characteristics of the basement membrane, epithelium, nuclei and lumen ([Fan](#)
150 [Fan et al., 2025](#)). In our pre-TraitHorizon analysis, non-negative matrix factorization ([F.](#)
151 [Fan et al., 2025](#); [Lee & Seung, 2001](#)) (rank = 14) followed by UMAP on these features
152 revealed 1,667 tubules forming outlier clusters. These were flagged in a separate “Needs
153 QC” column.
- 154 3. **“url” (optional)** – Allows linking each object to external resources such as secondary
155 visualizations, related analyses, or web-based viewers. In our use case, these URLs open
156 the corresponding tubule in the HistomicsUI viewer ([DigitalSlideArchive/HistomicsUI,](#)
157 [2025](#)) within the Digital Slide Archive (DSA) ([Gutman et al., 2017](#)), which displays the
158 tubule’s location within the original tissue slide for added histologic context.

159 4.2. Running TraitHorizon

160 The folder containing all image files should be provided to TraitHorizon’s command-line
161 interface via the `assets_path` argument. All images must reside in this directory, which
162 TraitHorizon uses to locate and display them in the user interface. Similarly, the path to the
163 TSV file should be specified using the `tsv_path` argument.

164 After starting the TraitHorizon server, a local URL (e.g., `http://localhost:5000`) will appear in
165 the terminal. Opening this URL in a web browser launches the TraitHorizon user interface
166 ([Figure 1](#)), allowing users to interactively explore images and their associated features.




Figure 1: TraitHorizon web application interface. (A) The interactive parallel coordinates plot displays 99 pathomic features from 260,201 tubules, wherein each blue line represents one tubule feature vector. (B) A violin plot shows the value distribution of a selected feature when hovered over. (C) The data grid presents each segmented tubular image, with the original PAS-stained tubule on the left and the segmentation overlay on the right (green: tubular basement membrane; white: tubular nuclei; black: tubular lumen; red: tubular epithelium). Associated feature data loaded from the TSV (e.g., filename and corresponding feature values) are also displayed. (D) The bottom status bar shows the number of instances currently visible. (E) The drop zone and export button allow users to save filter configurations and filtered row IDs. TraitHorizon allows users to customize which columns are included or excluded in the parallel coordinates plot. Here, the URL and filename columns were excluded from the parallel coordinates plot but are shown in the data grid.

4.3. Initial Exploration with Parallel Coordinates Plots and Data Grid

Initial exploration began with the parallel coordinates plot, which provided a global view of feature distributions across hundreds of thousands of tubules (Figure 1-A). When hovering over a feature axis, a violin plot is displayed (Figure 1-B), enabling distribution visualization and noting of properties such as skewness and normality. Here, when hovering over the feature TBM_THICK_MAX (maximum tubular basement membrane thickness), the violin plot revealed a strongly skewed distribution. This suggested the presence of tubules with extremal values, which required closer examination in the data grid (Figure 1-C) for qualitative assessment. Together, these visualizations help users quickly identify whether features are broadly consistent, contain outliers, or show irregular patterns that may reflect artifacts or biologically distinct characteristics.

4.4. Filtering Based on Feature(s)

TraitHorizon's interactive filtering feature allows users to isolate and inspect data subpopulation based on feature values. For example, we filtered tubules with high TBM_THICK_MAX values (maximum tubular basement membrane thickness) by dragging a brush (Yun, 2025) along the corresponding axis in the parallel coordinate plot (Figure 2 – green box). The filtering action highlighted the relevant instances (tubules) in red and displayed their rows in the data grid, where we observed that these tubules consistently showed thickened basement membranes (large green segments), a hallmark of atrophic tubules.




Figure 2: Interactive filtering with only one feature. TraitHorizon allows users to interactively filter instances based on any feature shown in the parallel coordinates plot. In this example, we hovered over the axis for the feature `TBM_THICK_MAX` (maximum basement membrane thickness) and dragged a brush to isolate tubules with values in the range from 30 to 40. In the parallel coordinate plot, 119 instances were dynamically highlighted in red after filtering (with this number indicated in the status bar), while all other ones remained visible but are greyed out. This color contrast makes it easier to visually compare the selected subset with the rest of the dataset across all other features. As shown here, the selected tubules were associated with visibly thickened basement membranes (enlarged green segments). Such interactive visualization enables users to explore and validate morphological patterns interactively. When hovering over a feature axis, a violin plot is generated for the selected subset (in red), providing an overview of feature-value distributions. In the data grid, only the filtered instances are shown.




Figure 3: TraitHorizon enables the identification of multi-feature expression patterns. To inspect a specific instance, users can first (A) hover over the instance in the data grid, which highlights the corresponding row in both the parallel coordinates plot and the data grid in blue. By visually examining its feature vector, users can identify features of interest. In this example, the selected tubule showed high expression of `TBM_THICK_MAX`, and (B) moderate expression of `LUMEN_THICK_MAX` (maximum thickness of the tubular lumen).

TraitHorizon's data grid also displays images alongside their respective quantitative features, allowing users to explore individual instances in greater detail without switching contexts to perform image lookups. Using the data grid we visually examined the subset of tubules with high TBM_THICK_MAX values and found an unexpected phenotype characterized by enlarged lumina. When hovering over these rows in the data grid, the corresponding feature vectors were highlighted in blue within the parallel coordinates plot (Figure 3-A). Comparing the qualitative images with corresponding feature vectors suggested the existence of a small subpopulation of tubules exhibiting both thickened TBM and lumen dilation (moderately elevated values of LUMEN_THICK_MAX, the maximum lumen thickness; Figure 3-B). To further study this subpopulation, we applied an additional brush on LUMEN_THICK_MAX in the parallel coordinates plot (Figure 4-2). Subsequent review with study pathologists indicated that these tubules represent a composite morphological phenotype, combining basement membrane thickening with an enlarged lumen (Fogo et al., 2016).

In addition, TraitHorizon allows users to export and import filters for reproducibility and record-keeping. Specifically, clicking the "Export Filter Settings" button results in the filtered row IDs and corresponding filter values being saved as a JSON file (Figure 4-3). Dragging the JSON file into the drop zone (Figure 4-4) results in the brush being automatically restored and the corresponding subset of objects being reshown in the data grid.




Figure 4: Interactive joint filtering with multiple features. TraitHorizon supports joint filtering across multiple features to identify more complex patterns. In this example, we first (1) dragged a brush to isolate tubules with TBM_THICK_MAX values between 30 and 40, and then (2) applied a second brush on LUMEN_THICK_MAX, selecting values between approximately 15 and 25. This combined filtering enabled the identification of tubules showing both basement membrane thickening and enlarged lumen. Users can export filter configurations and filtered row IDs as a JSON file by clicking the (3) 'Export Filter Settings' button. For reproducibility, the exported JSON can later be reloaded by dragging it into the (4) drop zone, automatically restoring the filtering criteria.

Overall, the seamless integration of feature visualization, interactive filtering, and image review allows users to (1) rapidly assess whether unusual feature values reflect histopathological patterns or segmentation/technical artifacts, (2) uncover novel biological insights (e.g., complex disease phenotypes) that might be missed using univariate methods, ultimately supporting more informed QC and pattern exploratory analysis, (3) enhance reproducibility by enabling the storage, sharing, and restoration of filtering configurations for downstream analyses.

4.5. External URL Integration for Contextual Visualization in HistomicsUI




Figure 5: TraitHorizon supports integration with external viewers such as HistomicsUI through user-defined URL links. In this example, the link to the DSA has been pre-configured so that clicking the “url” cell of the row (1) opened an external browser window displaying the target tubule centered in the DSA’s HistomicsUI viewer (2). This functionality can also be used to connect objects to other external resources, such as clinical data, supplementary analyses, or metadata information.

TraitHorizon supports the integration of external URLs for each instance, enabling users to link directly to additional resources such as image viewers, annotation tools or clinical databases, depending on the needs of the workflow (Figure 5). In our implementation, this feature was used to connect TraitHorizon with DSA. Clicking the URL associated with each tubule (Figure 5-1) opened the corresponding region in the DSA’s viewer (named HistomicsUI) at the WSI level (Figure 5-2), with the tubule centered in view. This allowed pathologists to assess the broader tubule microenvironment—including adjacent tubules, interstitial areas, and glomeruli—rather than solely examining isolated objects. Notably, this feature is not limited to external image viewers but can be extended to any external resource to support different analytical workflows.

4.6. Searching by Exact Value




Figure 6: Flexible feature search. Using the user-defined “Needs QC” indicator (where ‘1’ marks tubules flagged for quality control), we (1) clicked the magnifier icon, and (2) manually entered the flag value to search the dataset. (3) The subset highlighted tubules lacking lumens, as evidenced by low values across lumen-related features such as LUMEN_THICK_AVE/MIN/MAX/STD (average, minimum, maximum, and standard deviation of tubular lumen thickness). This search functionality enables rapid identification of instances with specific feature values that need further review.

Beyond numerical range filtering, TraitHorizon also supports searching by specifying exact values for selected features. Users can (1) click the magnifier icon (Figure 6-1), (2) select the desired feature, and (3) input a specific category or value to search the dataset (Figure 6-2). Using this functionality, we searched for the “Needs QC” indicator to investigate tubules flagged for review. Visual inspection revealed that many flagged tubules lacked a visible lumen, and quantitative analyses (via the parallel coordinates plot) confirmed consistently low values for lumen-related features such as LUMEN_THICK_AVE (average lumen thickness, Figure 6-3). TraitHorizon thus enables rapid identification and efficient validation of instances with specific feature values—something that is often cumbersome in workflows using traditional/static tools.

Importantly, this search capability is not restricted to quality control indicators. It can be applied to any features in the dataset, enabling users to flexibly isolate and analyze subsets of interest based on labels, classifications, or other features. This is particularly useful for features containing numerous unique categories or identifiers—for example, identifying a specific tubule by its filename or retrieving all tubules originating from the same patient.

236 4.7. Sorting Numerical Features




Figure 7: Sorting numerical features. TraitHorizon allows users to sort numerical features directly within the data grid by clicking on a column header. Features can be sorted in either descending (left) or ascending (right) order. In this example, the feature TBM_AREA (tubular basement membrane area) was sorted to facilitate comparison between tubules with large/thickened basement membranes and those with small/thin ones, helping users visually assess underlying morphological differences. In the descending view (left panel), the top-ranked tubules exhibited larger or thickened TBM regions (extensive green segments). Conversely, in ascending view (right panel), the top entries corresponded to tubules with smaller or thinner basement membranes. By sorting features interactively, users can identify qualitative gradients across ordered set of objects, revealing progressive feature changes within the dataset.

237 TraitHorizon also supports sorting of numerical features directly within the data grid by clicking
238 on a column header (Figure 7). In our workflow, we sorted tubules by TBM_AREA (tubular
239 basement membrane area) to explore morphological variation. Scrolling through the sorted grid
240 revealed a clear gradient in morphology—specifically, progressive thickening of the basement
241 membrane—enabling intuitive observation of trends from normal to atrophic tubules.

242 5. Discussion and Conclusions

243 TraitHorizon addresses a critical gap in image-based data analysis by enabling fast, interactive,
244 and image-linked exploration of high dimensional feature datasets. Traditional visualization
245 approaches often fail to couple image data with derived features and support interactive
246 plotting at scale. In contrast, TraitHorizon links each source image to its high dimensional
247 feature vector, supports real-time interactive filtering and sorting, and utilizes efficient rendering
248 strategies to scale to datasets with hundreds of thousands of images.

249 TraitHorizon's features empower researchers to discover meaningful patterns, perform quality
250 control, and generate insights for large object-feature datasets. Furthermore, TraitHorizon's
251 support for arbitrary features, multiple feature types, and external URLs enable its broad use
252 across multiple computational imaging fields.

253 TraitHorizon is available as a python package, and a Docker image is released for easy
254 installation (<https://traithorizon.com/>).

255 6. Acknowledgements

256 Research reported in this publication was supported by:

- (1) the National Institutes of Health (NIH) under the following awards: R01LM013864 and R01DK118431;
- (2) The Nephrotic Syndrome Study Network (NEPTUNE) is part of the Rare Diseases Clinical Research Network (RDCRN), which is funded by the NIH and led by the National Center for Advancing Translational Sciences (NCATS) through its Division of Rare Diseases Research Innovation (DRDRI). NEPTUNE is funded under grant number U54DK083912 as a collaboration between NCATS and the National Institute of Diabetes and Digestive and Kidney Diseases (NIDDK). Additional funding and/or programmatic support is provided by the University of Michigan, NephCure Kidney International, Alport Syndrome Foundation, and the Halpin Foundation. RDCRN consortia are supported by the RDCRN Data Management and Coordinating Center (DMCC), funded by NCATS and the National Institute of Neurological Disorders and Stroke (NINDS) under U2CTR002818;
- (3) Additional support was also provided by NephCure and the Henry E. Haller, Jr. Foundation;
- (4) Funding for the CureGN consortium is provided by U24DK100845, U01DK100846, U01DK100876, U01DK100866, and U01DK100867 from the National Institute of Diabetes and Digestive and Kidney Diseases (NIDDK). Patient recruitment is supported by NephCure. Dates of funding for first phase of CureGN was 9/16/2013-5/31/2019. Dates of funding for the second phase of CureGN was 6/1/2019 - 5/31/2024. Date for the fundings for the third phase of CureGN are 6/1/2024-5/31/2029

7. References

- Ambekar, A., Roohian, M., Liu, Q., Wang, B., Fan, F., Cassol, C., Lafata, K., Holzman, L., Mariani, L., Hodgins, J., Zee, J., Janowczyk, A., & Barisoni, L. (2025). *Glomerular Segmentation, Classification, and Pathomic Feature-based Prediction of Clinical Outcomes in Minimal Change Disease and Focal Segmental Glomerulosclerosis*. medRxiv. <https://doi.org/10.1101/2025.10.01.25336172>
- Automated Computational Detection of Interstitial Fibrosis, Tubular Atrophy, and Glomerulosclerosis | American Society of Nephrology*. (n.d.). Retrieved August 29, 2022, from <https://jasn.asnjournals.org/content/32/4/837.abstract>
- Barisoni, L., Nast, C. C., Jennette, J. C., & others. (2013). Digital pathology evaluation in the multicenter Nephrotic Syndrome Study Network (NEPTUNE). *Clinical Journal of the American Society of Nephrology (CJASN)*, 8(8), 1449–1459. <https://doi.org/10.2215/CJN.08370812>
- Bertucci, D., Hamid, M. M., Anand, Y., Ruangrotsakun, A., Tabatabai, D., Perez, M., & Kahng, M. (2022). *DendroMap: Visual Exploration of Large-Scale Image Datasets for Machine Learning with Treemaps*. arXiv. <https://doi.org/10.48550/arXiv.2205.06935>
- Chen, Yijiang, Wang, B., Demeke, D., Fan, F., Berthier, C., Mariani, L., Lafata, K., Holzman, L., Hodgins, J., Janowczyk, A., Barisoni, L., & Madabhushi, A. (2025). Clinical relevance of computational pathology analysis of interplay between kidney microvasculature and interstitial microenvironment. *Clinical Journal of the American Society of Nephrology*, 20(2), 239–255. <https://doi.org/10.2215/CJN.0000000597>
- Chen, Y., Zee, J., Janowczyk, A. R., Rubin, J., Toro, P., Lafata, K. J., Mariani, L. H., Holzman, L. B., Hodgins, J. B., Madabhushi, A., & Barisoni, L. (2023). Clinical relevance of computationally derived attributes of peritubular capillaries from kidney biopsies. *Kidney360*, 4(5), 648–658. <https://doi.org/10.34067/KID.0000000000000116>
- CureGN Study Rationale, Design, and Methods: Establishing a Large Prospective Observational Study of Glomerular Disease - PubMed*. (n.d.). Retrieved May 30, 2024, from <https://pubmed.ncbi.nlm.nih.gov/30420158/>

- 306 *D3 by Observable* | *The JavaScript library for bespoke data visualization*. (n.d.). Retrieved
307 July 9, 2025, from <https://d3js.org/>
- 308 *DigitalSlideArchive/HistomicsUI*. (2025). Digital Slide Archive. [https://github.com/](https://github.com/DigitalSlideArchive/HistomicsUI)
309 [DigitalSlideArchive/HistomicsUI](https://github.com/DigitalSlideArchive/HistomicsUI)
- 310 Echle, A., Rindtorff, N. T., Brinker, T. J., Luedde, T., Pearson, A. T., & Kather, J. N. (2021).
311 Deep learning in cancer pathology: A new generation of clinical biomarkers. *Br J Cancer*,
312 124(4), 686–696. <https://doi.org/10.1038/s41416-020-01122-x>
- 313 Fan, F., Liu, Q., Zee, J., & others. (2024). Computationally derived characterization of tubular
314 changes in relation to the development of interstitial fibrosis: FR-PO933. *Journal of the*
315 *American Society of Nephrology*, 35(10S). <https://doi.org/10.1681/ASN.20248v5y8bwn>
- 316 Fan, Fan, Liu, Q., Zee, J., Ozeki, T., Demeke, D., Yang, Y., Bitzer, M., O'Connor, C. L.,
317 Farris, A. B., Wang, B., Shah, M., Jacobs, J., Mariani, L., Lafata, K., Rubin, J., Chen, Y.,
318 Holzman, L., Hodgins, J. B., Madabhushi, A., ... Janowczyk, A. (2025). Clinical relevance of
319 computationally derived tubular features and their spatial relationships with the interstitial
320 microenvironment in minimal change disease/focal segmental glomerulosclerosis. *Kidney*
321 *International*, 0(0). <https://doi.org/10.1016/j.kint.2025.04.026>
- 322 Fan, Fan, Martinez, G., DeSilvio, T., Shin, J., Chen, Y., Jacobs, J., Wang, B., Ozeki, T.,
323 Lafarge, M. W., Koelzer, V. H., Barisoni, L., Madabhushi, A., Viswanath, S. E., &
324 Janowczyk, A. (2024). CohortFinder: An open-source tool for data-driven partitioning
325 of digital pathology and imaging cohorts to yield robust machine-learning models. *Npj*
326 *Imaging*, 2(1), 15. <https://doi.org/10.1038/s44303-024-00018-2>
- 327 Fan, F., Smith, C., Nair, V., & others. (2025). Integration of pathomic signatures and spatial
328 transcriptomics reveals tubular morphometric trajectories and enables patient stratification
329 in minimal change disease (MCD)/FSGS: TH-OR044. *Journal of the American Society of*
330 *Nephrology*, 36(10S). <https://doi.org/10.1681/ASN.2025ca1haevy>
- 331 Fogo, A. B., Lusco, M. A., Najafian, B., & Alpers, C. E. (2016). AJKD atlas of renal
332 pathology: Ischemic acute tubular injury. *American Journal of Kidney Diseases*, 67(5), e25.
333 <https://doi.org/10.1053/j.ajkd.2016.03.003>
- 334 Gadegbeku, C. A., Gipson, D. S., Holzman, L. B., Ojo, A. O., Song, P. X. K., Barisoni, L.,
335 Sampson, M. G., Kopp, J. B., Lemley, K. V., Nelson, P. J., Lienczewski, C. C., Adler, S.
336 G., Appel, G. B., Cattran, D. C., Choi, M. J., Contreras, G., Dell, K. M., Fervenza, F. C.,
337 Gibson, K. L., ... Kretzler, M. (2013). Design of the Nephrotic Syndrome Study Network
338 (NEPTUNE) to evaluate primary glomerular nephropathy by a multidisciplinary approach.
339 *Kidney Int*, 83(4), 749–756. <https://doi.org/10.1038/ki.2012.428>
- 340 Gutman, D. A., Khalilia, M., Lee, S., & others. (2017). The digital slide archive: A software
341 platform for management, integration, and analysis of histology for cancer research. *Cancer*
342 *Research*, 77(21), e75–e78. <https://doi.org/10.1158/0008-5472.CAN-17-0629>
- 343 Heinrich, J., & Weiskopf, D. (2013). State of the art of parallel coordinates. In M. Sbert &
344 L. Szirmay-Kalos (Eds.), *Eurographics 2013 - state of the art reports*. The Eurographics
345 Association. <https://doi.org/10.2312/conf/EG2013/stars/095-116>
- 346 Image file type and format guide - Media | MDN. (2025). In *MDN Web Docs*. [https://](https://developer.mozilla.org/en-US/docs/Web/Media/Guides/Formats/Image_types)
347 developer.mozilla.org/en-US/docs/Web/Media/Guides/Formats/Image_types
- 348 *Installation — HoloViews v1.21.0*. (n.d.). Retrieved July 3, 2025, from <https://holoviews.org/>
- 349 Janowczyk, A., Zuo, R., Gilmore, H., Feldman, M., & Madabhushi, A. (2019). HistoQC: An
350 Open-Source Quality Control Tool for Digital Pathology Slides. *JCO Clin Cancer Inform*,
351 3, CCI.18.00157. <https://doi.org/10.1200/CCI.18.00157>
- 352 Lee, D., & Seung, H. (2001). Algorithms for Non-negative Matrix Factorization. *Adv. Neural*
353 *Inform. Process. Syst.*, 13.

- 354 Maaten, L. van der, & Hinton, G. (2008). Visualizing Data using t-SNE. *Journal of Machine*
355 *Learning Research*, 9(86), 2579–2605. <http://jmlr.org/papers/v9/vandermaaten08a.html>
- 356 McInnes, L., Healy, J., & Melville, J. (2020). UMAP: Uniform Manifold Approximation and
357 *Projection for Dimension Reduction*. arXiv. <https://doi.org/10.48550/arXiv.1802.03426>
- 358 *Pallets/flask: The Python micro framework for building web applications*. (n.d.). Retrieved
359 October 28, 2025, from <https://github.com/pallets/flask/?tab=readme-ov-file>
- 360 *SlickGrid Home*. (n.d.). Retrieved July 9, 2025, from <https://slickgrid.net/Index.html>
- 361 TensorFlow Team. (2025). *Visualizing data using the embedding projector in TensorBoard*.
362 https://www.tensorflow.org/tensorboard/tensorboard_projector_plugin.
- 363 Yun, X. (2025). *BigFatDog/parcoords-es*. <https://github.com/BigFatDog/parcoords-es>

DRAFT